

YASH HOODA

Data Engineer · AI Engineer · Software Engineer

346-698-6864 | yash.hooda6@gmail.com | Richmond, TX | US Citizen

linkedin.com/in/yash-hooda-384430242 | github.com/yashhooda1 | yashhooda.ai | hoodahiring.ai

PROFESSIONAL SUMMARY

Data Engineer and AI Engineer with hands-on experience building Python, SQL, PySpark, REST API, Microsoft Fabric, Databricks, Azure SQL, Lakehouse, and Power BI solutions. Strong background in ETL/ELT development, data modeling, validation, dashboarding, and production analytics across franchise operations, healthcare supply chain, cybersecurity, and manufacturing datasets.

PYTHON & SQL HIGHLIGHTS

- Built Python and PySpark ETL/ELT pipelines using Pandas, NumPy, Requests, JSON parsing, Spark DataFrames, Delta Lake, and Microsoft Fabric notebooks.
- Wrote complex SQL queries using joins, CTEs, aggregations, window-style reporting logic, deduplication checks, reconciliation queries, and data quality validation scripts.
- Developed Python automation for REST API ingestion, incremental loads, pipeline monitoring, exception handling, data cleansing, anomaly detection, and BI-ready transformations.
- Worked with datasets ranging from thousands to 10M+ records across restaurant POS, healthcare operations, cybersecurity training, and manufacturing analytics.

TECHNICAL SKILLS

Programming: Python (Pandas, NumPy, PySpark, Requests), SQL (T-SQL, Spark SQL, Azure SQL), JavaScript, Java, C/C++, R

Data Engineering: ETL/ELT, REST API ingestion, JSON, PySpark, Databricks, Microsoft Fabric, Lakehouse Architecture, Delta Lake, Medallion Architecture, Data Modeling, Data Validation, Incremental Processing

Cloud/BI: Azure Synapse, Microsoft Fabric, Azure SQL, Power BI, DAX, PowerQuery M, Power Platform, GitHub Actions, Vercel

AI/ML: LangChain, OpenAI API, Claude API, RAG, Vector Stores, Embeddings, NLP, Scikit-Learn, TensorFlow, Prompt Engineering

WORK EXPERIENCE

GSR Group — Data Engineer

Houston, TX | March 2026 – Present

- Built an end-to-end Toast POS API ingestion pipeline for 25+ Nothing Bundt Cakes franchise locations using Python, SQL, Microsoft Fabric, Spark notebooks, and Bronze/Silver/Gold Lakehouse architecture.
- Developed Python-based REST API ingestion workflows to process restaurant sales, payments, refunds, voids, discounts, service charges, guests, and operational transaction data across multi-store franchise environments.
- Wrote complex SQL validation and reconciliation queries to identify duplicate order_guid and payment_guid records, validate financial reporting logic, and improve data accuracy for leadership dashboards.
- Engineered incremental ETL/ELT logic with control tables for automated backfills, scheduling, monitoring, batch processing, and data quality tracking across 25+ locations.
- Created centralized SQL data models and analytics layers powering downstream Power BI dashboards and Azure SQL/Fabric reporting systems for franchise leadership.
- Optimized PySpark and SQL transformations to handle Fabric compute constraints, scalable batching strategies, and efficient incremental data processing.

Ternio Group — Healthcare Supply Chain Analytics — Data Engineer (Contractor)

Sugar Land, TX | Sep 2025 – Feb 2026

- Engineered automated Python and PySpark ETL pipelines for multi-hospital surgical instrument, sterilization, and operating room datasets, producing clean Lakehouse tables for analytics.
- Developed advanced SQL queries, views, joins, and aggregations supporting anomaly detection, utilization reporting, compliance risk analysis, and executive Power BI dashboards.
- Automated data cleaning, standardization, and validation workflows using Python, reducing manual reporting workload by 70%.
- Built scalable Microsoft Fabric dataflows integrating SQL, PowerQuery M, and Spark notebooks for cloud-based healthcare operations analytics.
- Optimized SQL queries and Lakehouse structures across high-volume datasets exceeding 10M rows.

Archrock — Natural Gas Compression Services — Data Engineer Intern

Houston, TX | May 2025 – Aug 2025

- Built Python scripts to clean, transform, and analyze cybersecurity training datasets for phishing simulation reporting and compliance tracking.
- Developed SQL-based reporting logic and data quality checks to evaluate phishing click rates, employee training completion, and campaign effectiveness.
- Used Microsoft Fabric, Lakehouses, and Azure Synapse pipelines to process unstructured cybersecurity training data into clean analytical tables.
- Automated recurring Power BI reporting workflows, reducing reporting delays by 50%.

Century Instrument & Machine Company (CIMCO) — Data Analyst (Contractor)

Houston, TX | Feb 2025 – May 2025

- Used SQL and Excel-based analysis to evaluate supplier pricing, delivery times, quality metrics, and operational performance trends, identifying 10% potential cost savings.
- Cleaned and analyzed 1,000+ rows of operational data, built automated reports and dashboards, and reduced reporting time by 30%.
- Created training and operations documentation while evaluating cost-effective software solutions that reduced projected training expenses by 25%.

Outlier.AI — AI Trainer (Freelance)

Remote | May 2024 – Feb 2025

- Evaluated AI model outputs, chatbot responses, and Python/NumPy code for correctness, reliability, and instruction-following quality.
- Reviewed NLP and coding tasks involving prompt engineering, model behavior, and output quality assessment.

Parijat Controlware — Software Engineer Intern

Houston, TX | May 2023 – Jun 2023

- Developed PLC-based automation logic in a team environment for industrial control systems.
- Applied object detection and machine learning using RoboFlow and Python for industrial automation use cases.

AI & PERSONAL PROJECTS

Personal Portfolio Website — RAG Chatbot, AI Agents, Strava Analytics

Claude/OpenAI APIs, RAG, Upstash Vector, Redis, Node.js, Vercel | 2025 – Present

- Architected a full-stack AI portfolio site featuring a security-focused chatbot gateway, multi-model support, and RAG over resume/project content via Upstash Vector.
- Built Strava Intelligence analytics with CTL/ATL/Form tracking, pace zones, AI coaching insights, per-run weather data, and race prediction logic.
- Shipped live flight tracker, live sports scores, OpenAI TTS voice conversation, memory-backed chat features, and network analyzer tooling.
- Deployed on Vercel with serverless Node.js backend, Upstash Redis cross-session memory, and SEO optimization.

Apex CRM — AI-Powered Sales CRM

React, Claude Sonnet API, Kanban, Lead Scoring | 2025

- Built an AI-powered React sales CRM featuring Kanban pipeline tracking, lead scoring, automated email drafting, and close strategy generation using Claude Sonnet.
- Packaged with GitHub documentation and portfolio framing around data engineering patterns, AI integration, and production-style software delivery.

HoodaAgents — AI Agent Framework

Python, LangChain, OpenAI API, Streamlit, Tavily Search, RAG | Jan 2025 – Jun 2025

- Designed and built a full-stack Python AI agent framework with dynamic memory recall, personalized chatbot personas, and RAG pipelines.
- Deployed public Streamlit demo with Tavily web search integration, Hindi language support, vector store memory, and modular tool workflows.

NBC Franchise MIS Dashboard

Microsoft Fabric, Azure SQL, Power BI, GitHub Actions | 2025

- Built a franchise management dashboard for 25 Nothing Bundt Cakes locations covering YoY sales comparisons, refunds, discounts, and store-level filters.
- Implemented automated daily GitHub Actions refresh pipeline from Fabric SQL warehouse via service principal; added Excel export and NBC brand theming.

Liver Cancer Prediction — ML Model

Python, Scikit-Learn, Pandas, NumPy, Matplotlib | Feb 2025

- Analyzed patient records dataset through exploratory data analysis; implemented Random Forest and Logistic Regression models with interpretable feature importance visualizations.

EDUCATION

University of Texas at Dallas — Bachelor of Science, Computer Science | Richardson, TX | May 2024

CERTIFICATIONS

- Databricks Certified Data Engineer Associate — Databricks | Apr 2026 – Apr 2028
- Microsoft Power Platform Fundamentals (PL-900) — Microsoft | Apr 2026
- IBM Professional AI Engineering Certificate — Coursera / IBM | Aug 2024
- IBM Professional Data Science Certification — Coursera / IBM | Nov 2024
- TensorFlow Developer Certification — DeepLearning.AI | Jan 2024
- Vanderbilt University AI Prompt Engineering Certification — Coursera / Vanderbilt | Feb 2025