

YASH HOODA

AI Engineer · Data Engineer · Software Engineer

346-698-6864 | yash.hooda6@gmail.com | Richmond, TX | US Citizen

linkedin.com/in/yash-hooda-384430242 | github.com/yashhooda1 | yashhooda.ai | hoodahiring.ai

PROFESSIONAL SUMMARY

Data and AI Engineer with 2+ years of experience building production-grade ETL pipelines, AI-powered applications, and cloud-native data platforms. Proven track record shipping end-to-end systems across franchise analytics, healthcare supply chain, and full-stack AI — from REST API ingestion to Power BI dashboards to multi-model LLM apps with RAG, agents, and real-time vector search. Databricks Certified. BS Computer Science, UT Dallas.

TECHNICAL SKILLS

Languages: Python (Pandas, NumPy, PySpark, Requests), SQL (T-SQL, Spark SQL, Azure SQL), JavaScript, Java, C/C++, R

Data Engineering: ETL/ELT, REST API Ingestion, Delta Lake, Medallion Architecture, Microsoft Fabric, Databricks, PySpark, Incremental Processing, Data Modeling

AI / ML: LangChain, OpenAI API, Claude API (Anthropic), RAG, Vector Stores (FAISS, Chroma, Upstash), Embeddings, NLP, Scikit-Learn, TensorFlow, Prompt Engineering, LLM Fine-Tuning

Cloud & BI: Azure (Synapse, Fabric, Foundry), Power BI, DAX, PowerQuery M, GitHub Actions, Vercel, Node.js Serverless

Tools: Git, GitHub, Postman, VS Code, Databricks, Azure SQL, Upstash Redis

CERTIFICATIONS

- **Databricks Certified Data Engineer Associate** — Databricks | Apr 2026 – Apr 2028
- **IBM Professional AI Engineering Certificate** — Coursera / IBM | Aug 2024
- **TensorFlow Developer Certification** — DeepLearning.AI | Jan 2024
- **IBM Professional Data Science Certification** — Coursera / IBM | Nov 2024
- **Microsoft Power Platform Fundamentals (PL-900)** — Microsoft | Apr 2026
- **Vanderbilt University AI Prompt Engineering** — Coursera / Vanderbilt | Feb 2025

WORK EXPERIENCE

GSR Group — Data Engineer

Houston, TX | Mar 2026 – Present

- Built an end-to-end Toast POS API ingestion pipeline for 25+ Nothing Bundt Cakes franchise locations using Python, PySpark, Microsoft Fabric, and Bronze/Silver/Gold Lakehouse architecture — powering dashboards tracking millions in franchise revenue.
- Engineered incremental ETL/ELT workflows with control tables for automated backfills, scheduling, monitoring, and data quality tracking across restaurant sales, payments, refunds, voids, discounts, and service charges.
- Wrote complex SQL validation and reconciliation queries to detect duplicate order_guid and payment_guid records, enforce financial reporting accuracy, and surface data quality issues for leadership review.
- Created centralized SQL data models and analytics layers powering downstream Power BI dashboards and Azure SQL reporting for franchise leadership across 25+ store locations.
- Optimized PySpark and SQL transformations to handle Fabric compute constraints, scalable batching strategies, and efficient large-scale incremental data processing.

Ternio Group — Healthcare Supply Chain Analytics — Data Engineer (Contract)

Sugar Land, TX | Sep 2025 – Feb 2026

- Engineered automated Python and PySpark ETL pipelines for multi-hospital surgical instrument, sterilization, and OR datasets — reducing manual reporting workload by 70% and replacing legacy Excel processes with cloud-native automation.
- Built scalable Microsoft Fabric dataflows integrating SQL, PowerQuery M, and Spark notebooks for anomaly detection, OR utilization analysis, and compliance risk monitoring across 10M+ row datasets.
- Developed advanced SQL queries, views, and aggregations powering executive Power BI dashboards used by hospital leadership to monitor throughput, utilization, and multi-million-dollar compliance metrics.

Archrock — Natural Gas Compression Services — Data Engineer (Contract)

Houston, TX | May 2025 – Aug 2025

- Built Python and SQL pipelines to clean, transform, and analyze cybersecurity training datasets for phishing simulation tracking and compliance reporting; reduced reporting delays by 50%.
- Used Microsoft Fabric, Lakehouses, and Azure Synapse to process unstructured cybersecurity training data into clean analytical tables; automated recurring Power BI reporting workflows.

Century Instrument & Machine Company (CIMCO) — Data Analyst (Contract) Houston, TX | Feb 2025 – May 2025

- Used SQL and Excel analysis to evaluate supplier pricing, delivery times, and quality metrics — identifying 10% potential cost savings; built automated dashboards that cut reporting time by 30%.

Outlier.AI — AI/ML Data Specialist (Freelance) Remote | May 2024 – Feb 2025

- Evaluated AI model outputs, chatbot responses, and Python/NumPy code for correctness and quality; contributed feedback that drove 90%+ improvement in model performance across NLP and coding tasks.

Parijat Controlware — Software Engineer Intern Houston, TX | May 2023 – Jun 2023

- Developed PLC-based automation logic for industrial control systems; applied object detection and ML using RoboFlow and Python for industrial automation use cases.

AI & PERSONAL PROJECTS

yashhooda.ai — Personal AI Portfolio Platform | Node.js, Claude/OpenAI APIs, RAG, Upstash Vector/Redis, Vercel | 2025 – Present

- Architected a production full-stack AI platform with 7-layer security gateway (per-IP rate limiting, jailbreak detection, Redis kill switch, CIDR blocks), multi-model support (Claude, GPT-4, Grok, Gemini, Llama), and RAG over resume/project content via Upstash Vector.
- Built Strava Intelligence analytics with CTL/ATL/Form tracking, pace zones, AI coaching insights, per-run weather data, and marathon race prediction logic.
- Shipped live flight tracker (FlightAware + ADS-B), sports scores (ESPN API), OpenAI TTS voice conversation, coding agent with agentic loop, and autonomous research agent with Vercel Cron.

HoodaHiring.ai — AI Hiring Platform | Python, Railway, Docker, Custom Domain | 2025 – Present

- Built and deployed a live AI hiring platform on Railway with custom domain (hoodahiring.ai), Dockerfile-based containerization, and Google Search Console SEO verification.

Apex CRM — AI-Powered Sales CRM | React, Claude Sonnet API, Kanban, Lead Scoring | 2025

- Built a full-stack AI sales CRM featuring Kanban pipeline, intelligent lead scoring, automated email drafting, and close strategy generation using Claude Sonnet — packaged with complete GitHub documentation.

HoodaAgents — AI Agent Framework | Python, LangChain, OpenAI API, Streamlit, Tavily Search, RAG | Jan 2025 – Jun 2025

- Designed and deployed a full-stack AI agent framework with dynamic memory recall, personalized personas, RAG pipelines, Tavily web search, Hindi language support, and vector store memory.

NBC Franchise MIS Dashboard | Microsoft Fabric, Azure SQL, Power BI, GitHub Actions | 2025

- Built a franchise management dashboard for 25 Nothing Bundt Cakes locations (YoY sales, refunds, discounts, store-level filters) with daily GitHub Actions refresh pipeline via service principal and Excel export.

Databricks LLM Fine-Tuning Pipeline | Mistral-7B, Llama 3 8B, QLoRA, MLflow, HuggingFace Hub | 2025

- Built end-to-end fine-tuning pipeline for a PySpark coding assistant using QLoRA on Mistral-7B/Llama 3 8B, MLflow experiment tracking, Unity Catalog integration, and LoRA adapter deployment to HuggingFace Hub.

EDUCATION

University of Texas at Dallas — Bachelor of Science, Computer Science Richardson, TX | May 2024