

YASH HOODA

Data Engineer · AI Engineer

Richmond, TX | 346-698-6864 | yash.hooda6@gmail.com | US Citizen
linkedin.com/in/yash-hooda-384430242 | github.com/yashhooda1 | yashhooda.ai | hoodahiring.ai

SUMMARY

Data Engineer with experience building production ETL and AI systems on Azure, Databricks, and Microsoft Fabric — from incremental REST-API ingestion and medallion-architecture Delta Lake pipelines to executive Power BI reporting. Also ship production LLM applications — hybrid RAG, agentic workflows, and multi-model gateways — serving live traffic. Databricks Certified Data Engineer. BS Computer Science, UT Dallas.

TECHNICAL SKILLS

Languages Python (Pandas, NumPy, PySpark, Requests), SQL (T-SQL, Spark SQL), JavaScript/TypeScript, Java, C/C++, R
Data Engineering ETL/ELT, REST API Ingestion, Delta Lake, Medallion Architecture, Incremental Processing & Control Tables, Data Modeling, Data Quality & Reconciliation, Microsoft Fabric, Databricks
AI / ML RAG (hybrid retrieval, RRF, reranking), LLM Agents, LangChain/LangGraph, Claude & OpenAI APIs, Vector Stores (Upstash, FAISS, Chroma), QLoRA Fine-Tuning, MLflow, Evals, Scikit-Learn
Cloud & BI Azure (Synapse, Fabric, SQL), Power BI, DAX, PowerQuery M, GitHub Actions, Docker, Vercel, Node.js Serverless, Redis

EXPERIENCE

Ternio Group (Healthcare Supply Chain) — Data Engineer, Contract Sugar Land, TX | Sep 2025 – Feb 2026

- Automated multi-hospital ETL for surgical instrument, sterilization, and OR datasets in Python/PySpark — **cutting manual reporting workload 70%** and retiring a legacy Excel process.
- Built Microsoft Fabric dataflows across **10M+ row datasets** integrating SQL, PowerQuery M, and Spark notebooks for anomaly detection, OR utilization analysis, and compliance risk monitoring feeding executive Power BI dashboards.

Archrock (Natural Gas Compression) — Data Engineer Intern Houston, TX | May 2025 – Aug 2025

- Built Python and SQL pipelines processing unstructured cybersecurity training data into clean analytical tables via Fabric, Lakehouses, and Azure Synapse; automated recurring Power BI reporting and **reduced reporting delays 50%**.

Century Instrument & Machine (CIMCO) — Data Analyst, Contract Houston, TX | Feb 2025 – May 2025

- Analyzed supplier pricing, delivery, and quality metrics in SQL — identifying **10% potential cost savings** and building automated dashboards that **cut reporting time 30%**.

Outlier.AI — AI/ML Data Specialist Freelance, Remote | May 2024 – Feb 2025

- Evaluated LLM outputs and Python/NumPy code for correctness and instruction-following across NLP and coding tasks, contributing to RLHF training data.

SELECTED PROJECTS

NBC MIS Dashboard — Franchise Analytics Pipeline Python, Toast Orders API, GitHub Actions, Chart.js, Leaflet, Vercel

- Built an **end-to-end Toast POS ingestion pipeline** pulling the Orders API and aggregating per store-day into an idempotent cache — powering a live executive dashboard for **25+ franchise locations**: KPI strip, US geo map with drill-down, same-store YoY, and payment-mix attribution across Credit / DoorDash / Uber Eats / Grubhub / ezCater / Cash.
- Automated incremental refresh via **GitHub Actions with keyed replay** so any date window can be reprocessed without double-counting revenue; deployed on Vercel.

yashhooda.ai — Production AI Platform Node.js, Claude/OpenAI, LangGraph, Upstash Vector/Redis, Vercel

- Hybrid RAG pipeline**: dense + sparse retrieval with Reciprocal Rank Fusion, CRAG 1–5 relevance grading, conditional web-search fallback, and cross-encoder reranking — with an offline eval harness running in CI.
- 7-layer security gateway** (per-IP rate limiting, jailbreak detection, VPN/datacenter fingerprinting, content-guard auto-ban, Redis kill switch) and a 5-provider LLM fallback chain. Withstood a sustained real-world attack campaign in production.

PySpark Coding Assistant — LLM Fine-Tuning Pipeline Mistral-7B, QLoRA, PEFT, MLflow, HuggingFace

- Built a **synthetic training corpus via LLM distillation** from hand-curated production patterns, with an automated validator (AST parse, required-API check, deprecated-API and pandas rejection) — after a 122K-row public code dataset yielded only 58 usable PySpark examples.
- Fine-tuned Mistral-7B with QLoRA (4-bit NF4, rank-16 LoRA) on a T4; final eval loss 0.76, perplexity 1.69. **Adapter and dataset published to HuggingFace Hub** with a model card documenting measured accuracy and failure cases — corpus size identified as the binding constraint for v2.

hoodahiring.ai — LLM Resume Intelligence Python, Docker, Railway

- Live containerized application that parses resumes into structured candidate intelligence and scores fit against a job description; deployed on a custom domain.

EDUCATION & CERTIFICATIONS

University of Texas at Dallas — B.S. Computer Science | Richardson, TX | May 2024

Databricks Certified Data Engineer Associate (2026–2028) · Microsoft PL-900 (2026) · IBM AI Engineering Professional Certificate · TensorFlow Developer Certification