

YASH HOODA

AI Engineer · Data Engineer · Software Engineer

346-698-6864 | yash.hooda6@gmail.com | Richmond, TX | US Citizen

[linkedin.com/in/yash-hooda-384430242](https://www.linkedin.com/in/yash-hooda-384430242) | github.com/yashhooda1 | yashhooda.ai | hoodahiring.ai

EDUCATION

University of Texas at Dallas — Richardson, TX
Bachelor of Science, Computer Science

May 2024

WORK EXPERIENCE

GSR Group — Data Engineer

Houston, TX | March 2026 – Present

- ▶ Built an end-to-end Toast POS API ingestion pipeline for 25+ Nothing Bundt Cakes franchise locations using Microsoft Fabric, Spark notebooks, and medallion architecture (Bronze/Silver/Gold).
- ▶ Engineered incremental ETL/ELT workflows to process restaurant sales, payments, refunds, voids, discounts, and operational transaction data at scale across 25+ locations.
- ▶ Developed centralized data models and analytics pipelines powering downstream Power BI dashboards and Azure SQL reporting systems for franchise leadership.
- ▶ Designed control tables and orchestration logic for automated backfills, scheduling, monitoring, and data quality tracking across multi-store franchise environments.
- ▶ Optimized large-scale ingestion workflows to handle Fabric compute capacity constraints, batching strategies, and scalable incremental processing.
- ▶ Stack: REST APIs, JSON, PySpark, Delta Lake, SQL, Python, Azure cloud analytics infrastructure.

Ternio Group — Healthcare Supply Chain Analytics — Data Engineer (Contractor)

Sugar Land, TX | Sep 2025 – Feb 2026

- ▶ Engineered automated Python + PySpark ETL pipelines for multi-hospital surgical instrument and sterilization datasets, producing clean, production-grade Lakehouse tables.
- ▶ Built scalable Microsoft Fabric dataflows integrating SQL, PowerQuery M, and Spark notebooks for anomaly detection, OR utilization, and compliance risk analysis.
- ▶ Developed real-time Power BI dashboards used by hospital leadership to monitor throughput, utilization, and multi-million-dollar compliance metrics.
- ▶ Reduced manual reporting workload by 70% by replacing Excel and legacy processes with cloud-native automation pipelines.
- ▶ Optimized SQL queries and Lakehouse structures across high-volume datasets (10M+ rows).

Archrock — Natural Gas Compression Services — Data Engineer Intern

Houston, TX | May 2025 – Aug 2025

- ▶ Automated phishing simulation and employee training metrics tracking using Power BI, reducing reporting delays by 50%.
- ▶ Used Microsoft Fabric, Lakehouses, and Azure Synapse pipelines to process unstructured cybersecurity training data into clean, analyzable tables.
- ▶ Developed logic to evaluate phishing click rates, training effectiveness, and data quality issues using Python and SQL.

Century Instrument & Machine Company (CIMCO) — Data Analyst (Contractor)

Houston, TX | Feb 2025 – May 2025

- ▶ Conducted supplier performance analysis evaluating pricing, delivery times, and quality metrics, identifying 10% potential cost savings.
- ▶ Analyzed and cleaned 1,000+ rows of operational data; built automated reports and dashboards, reducing reporting time by 30%.
- ▶ Designed a comprehensive machinist training document; evaluated four cost-effective training software solutions, reducing training expenses by 25%.

Outlier.AI — AI Trainer (Freelance)

Remote | May 2024 – Feb 2025

- ▶ Evaluated AI model outputs, chatbot responses, and Python/NumPy code for accuracy and quality; feedback contributed to 90%+ improvement in model performance.
- ▶ Participated in NLP and AI prompt evaluation projects, developing deep expertise in LLM behavior, output quality assessment, and prompt engineering.

Parijat Controlware — Software Engineer Intern

Houston, TX | May 2023 – Jun 2023

- ▶ Developed PLC-based automation logic in a team environment for industrial control systems.
- ▶ Applied object detection and machine learning using RoboFlow and Python for industrial automation use cases.

AI & PERSONAL PROJECTS

Personal Portfolio Website — RAG Chatbot, AI Agents, Strava Analytics | Claude/OpenAI APIs, RAG, Upstash Vector, Redis, Node.js, Vercel 2025 – Present

- ▶ Architected a full-stack AI portfolio site featuring a 7-layer security gateway chatbot, multi-model support (Claude, GPT, Grok, Gemini, Llama), and RAG over resume/project content via Upstash Vector.
- ▶ Built Strava Intelligence analytics with CTL/ATL/Form tracking, pace zones, AI coaching insights, per-run weather data, and marathon race predictions.
- ▶ Shipped live flight tracker (FlightAware + ADS-B Exchange), live sports scores (ESPN API), OpenAI TTS voice conversation tab, and network analyzer tool.
- ▶ Deployed on Vercel with serverless Node.js backend, Upstash Redis cross-session memory, and Google Search Console/SEO optimization.

Apex CRM — AI-Powered Sales CRM | React, Claude Sonnet API, Kanban, Lead Scoring 2025

- ▶ Built a full AI-powered React sales CRM featuring Kanban pipeline, intelligent lead scoring, automated email drafting, and close strategy generation using Claude Sonnet.
- ▶ Packaged with complete GitHub documentation and portfolio framing around data engineering patterns and AI integration.

HoodaAgents — AI Agent Framework | Python, LangChain, OpenAI API, Streamlit, Tavily Search, RAG Jan 2025 – Jun 2025

- ▶ Designed and built a full-stack AI agent framework with dynamic memory recall, personalized chatbot personas, and RAG pipelines.
- ▶ Deployed public demo with Streamlit UI, Tavily web search integration, Hindi language support, and vector store memory.

NBC Franchise MIS Dashboard | Microsoft Fabric, Azure SQL, Power BI, GitHub Actions 2025

- ▶ Built a franchise management dashboard for 25 Nothing Bundt Cakes locations covering YoY sales comparisons, refunds, discounts, and store-level filters.
- ▶ Implemented automated daily GitHub Actions refresh pipeline from Fabric SQL warehouse via service principal; added Excel export and NBC brand theming.

Liver Cancer Prediction — ML Model | Python, Scikit-Learn, Pandas, NumPy, Matplotlib Feb 2025

- ▶ Analyzed patient records dataset via EDA; implemented and optimized Random Forest and Logistic Regression models with interpretable feature importance visualizations.

TECHNICAL SKILLS

AI / ML: LangChain, OpenAI API, Claude API (Anthropic), RAG, Vector Stores (FAISS, Chroma), Embeddings, NLP, TensorFlow, Scikit-Learn, Prompt Engineering, LLM Fine-Tuning

Data Engineering: PySpark, SQL, Microsoft Fabric, Databricks, ETL/ELT, Lakehouse Architecture (Bronze/Silver/Gold), Delta Lake, Medallion Architecture, REST API Ingestion

Programming: Python (Pandas, NumPy, PySpark), SQL, JavaScript, Java, C/C++, R

Cloud & Platforms: Azure (Synapse, Fabric, Foundry, Dataverse), Google Cloud, IBM Cloud, Microsoft Power Platform

BI & Visualization: Power BI, DAX, PowerQuery M, Streamlit, Jupyter Notebooks

Big Data: Apache Spark, Spark DataFrames, Delta Lake, Data Pipelines, Batch & Incremental Processing

Tools: Git, GitHub, GitHub Actions, VS Code, Postman, Databricks, Azure SQL

AI Tools: Claude (Anthropic), ChatGPT (OpenAI), Gemini, Meta Llama, Ollama, DeepSeek R-1, Grok-3, GitHub Copilot

CERTIFICATIONS

- **Databricks Certified Data Engineer Associate** — Databricks *Apr 2026 – Apr 2028*
- **Microsoft Power Platform Fundamentals (PL-900)** — Microsoft *April 2026*
- **IBM Professional AI Engineering Certificate** — Coursera / IBM *August 2024*
- **TensorFlow Developer Certification** — DeepLearning.AI *January 2024*
- **IBM Professional Data Science Certification** — Coursera / IBM *November 2024*
- **Vanderbilt University AI Prompt Engineering Certification** — Coursera / Vanderbilt *February 2025*